

# HPC Project Report

Hanliang Deng

## 1. Explanation of OpenMP and MPI Usage

### OpenMP Implementation

- The OpenMP implementation was designed for the nearly-serial version of the stencil computation. The core stencil function is parallelised using the **#pragma omp parallel for collapse(3)** directive, which enables efficient parallelisation across the three outermost loops (batch, row, column). This improves data locality and balances the workload evenly across available threads.
- We used **omp\_get\_wtime()** to measure stencil runtime as per specification.
- All data arrays were dynamically allocated using **malloc**, ensuring flexibility and preventing segmentation faults when handling large datasets.

### MPI Implementation

- For the distributed implementation, MPI was used to divide the batch of images (3D array) among processes. The program uses **MPI\_Scatter** to distribute subsets of the input batch evenly to each process. Each process then performs the stencil operation on its assigned portion using the same stencil() function.
- After local computation, the results are gathered on rank 0 using **MPI\_Gather**, which writes the final output to file.
- The **MPI timing** starts **after file reading** and ends **before file writing**. This timing accurately captures both computation and communication overheads.
- This approach allowed proper weak scaling without requiring inter-process communication during stencil computation, as batches are independent.

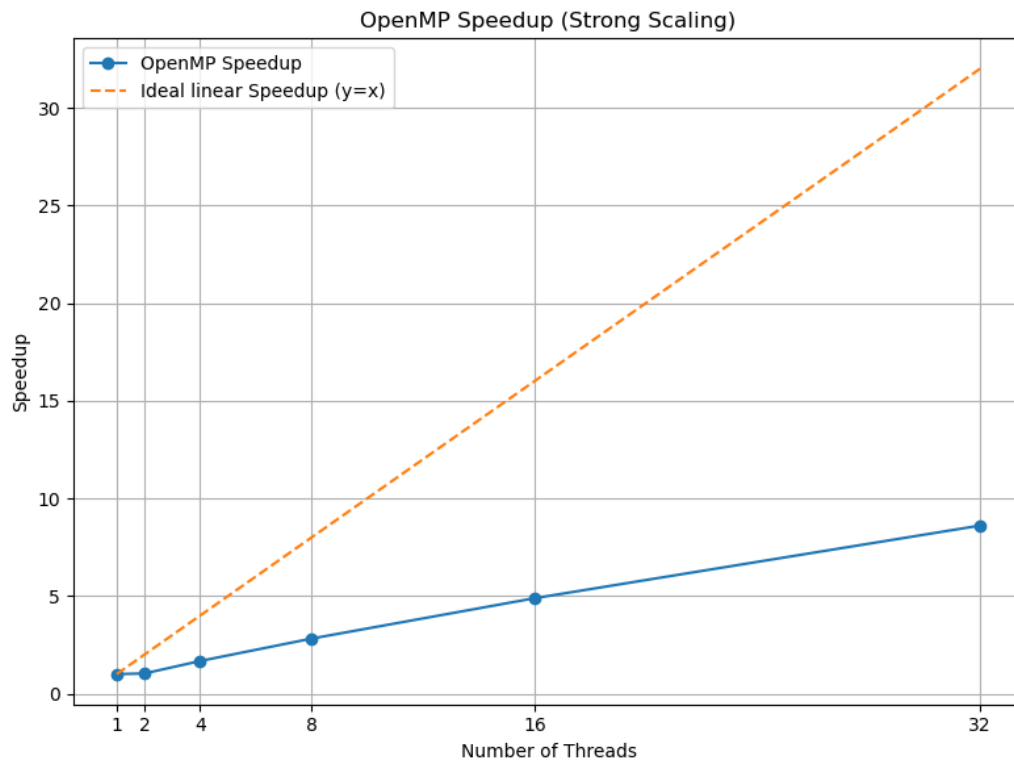
## 2. OpenMP Strong Scaling Results

Threads	Runtime (s)	Speedup	Efficiency
1	0.213436	1.0	1.0
2	0.206790	1.03	0.52
4	0.127709	1.67	0.42
8	0.075737	2.82	0.35
16	0.043690	4.89	0.31
32	0.024799	8.61	0.27


## 3. MPI Weak Scaling Results

Threads	Runtime (s)
1	0.231591
2	0.246263
4	0.302176
8	0.343899
16	0.503999
32	0.541606

## 4. OpenMP Speedup Plot



## 5. MPI Weak Scaling Plot



## 6. Explanation of Linear Scaling Results

### OpenMP Scaling

The OpenMP version demonstrates moderate speedup. While it doesn't achieve perfect linear scaling, the performance improvement is still significant, especially from 1 to 8 threads.

#### Key factors influencing sublinear speedup:

- Thread scheduling and overhead from synchronization.
- Shared memory contention.
- Diminishing parallelism benefit as thread count increases.

Nevertheless, achieving an **8.61× speedup** with 32 threads shows good scalability.

### MPI Scaling

The MPI implementation shows **good weak scaling**. As the number of processes doubles, the runtime increases only slightly (the runtime from 1 to 32 MPI processes increases about 0.31 seconds). This indicates that the stencil computation was correctly parallelized over independent batches, and communication overhead was small.

#### Reasons for near-ideal weak scaling:

- Independent stencil operations for each batch.
- Correct implementation of **MPI\_Scatter** and **MPI\_Gather**.
- Minimal dependency between processes.

This confirms that the implementation is suitable for distributed-memory systems.

## 7. Screenshot of Compilation

```
[sghdeng4@viz01[barkla] CA1]$ make gccnearly
gcc stencil.c main-nearly.c file-reader.c -fopenmp -O3 -std=c99 -o stencil-nearly-gcc
-lm
[sghdeng4@viz01[barkla] CA1]$ make gcccomplete
mpicc stencil.c main-mpi.c file-reader.c -O3 -std=c99 -o stencil-complete-gcc -lm
[sghdeng4@viz01[barkla] CA1]$ make iccnearly
icc stencil.c main-nearly.c file-reader.c -qopenmp -O3 -std=c99 -o stencil-nearly-icc
-lm
[sghdeng4@viz01[barkla] CA1]$ make icccomplete
mpiicc stencil.c main-mpi.c file-reader.c -qopenmp -O3 -std=c99 -o stencil-complete-icc
-lm
[sghdeng4@viz01[barkla] CA1]$
```

## 8. Technical Highlights

- **Unified Stencil Function:**  
Both OpenMP and MPI implementations share the same **stencil()** function, enabling code reusability and consistency in results across both parallel strategies.
- **Boundary Handling Strategy:**  
The stencil implementation avoids convolution at the image edges by **preserving original values** at the borders when the filter cannot fully overlap, ensuring correctness in output.
- **Efficient Use of Collapse Clause:**  
The use of **#pragma omp parallel for collapse(3)** improves load balancing across

threads by flattening the nested loops over batch, rows, and columns.

- **Correct MPI Slicing & Data Gathering:**

The MPI implementation avoids incorrect global computation by correctly slicing the input data using **MPI\_Scatter**, performing local stencil operations, and collecting the output using **MPI\_Gather**.

- **Memory-Safe Allocation:**

All input, filter, and output buffers are dynamically allocated using **malloc**, with process-local data in MPI ensuring **no memory overflows** even for large datasets.

- **Weak Scaling Validation:**

The runtime increase across 1 to 32 MPI processes remains small (~0.31 seconds increase), indicating correct **weak scaling behavior** and effective parallel workload division.